Aligning Open-source LLMs Using Reinforcement Learning + Feedback



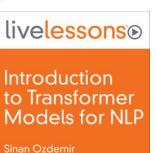
/sinanuozdemir/odsc-llm-alignment



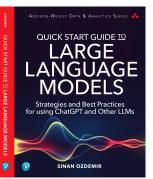
Sinan OzdemirData Scientist, Entrepreneur,
Author, Lecturer



Welcome!



video





My name is **Sinan Ozdemir** (in/sinan-ozdemir + @prof oz)

- Current founder of Loop Genius (using AI to help entrepreneurs get their first 100 customers)
- Current **lecturer** for O'Reilly and Pearson
- Founder of Kylie.ai (Funded by OpenAl Founder + Acquired)
- Masters in Theoretical Math from Johns Hopkins
- Former lecturer of Data Science at Johns Hopkins

Author of ML textbooks and online series, including

- The Principles of Data Science
- Introduction to Transformer Models for NLP
- **Quick Start Guide to LLMs**



On the Docket

- A (very short) primer on Alignment & Reinforcement Learning from Feedback
- Example Aligning FLAN-T5's summaries to be more neutral and grammatically correct
- Case Study Instruction Aligning an open-source LLM

- Evaluating Alignment



Alignment vs Reinforcement Learning

- **Alignment** is training a model to produce outputs more in line with human preference and expectation
- **Reinforcement Learning (RL)** is a machine learning technique that uses sequential feedback to teach an "agent" how to behave in an "environment".

RL is the most talked about method of alignment but not the only option! OpenAI popularized the method in 2022 specifically using RLHF (RL from Human Feedback)



Alignment Makes LLMs Do What We Want

Alignment in LLMs – Refers to how a language model understands and responds to input prompts in a way that aligns with the user's expectations. Humans (or AI) in the loop judge and reward LLM outputs to ensure that the model's responses are "in line with" what the user intended or expected.

Is the Earth flat?

Yes.

GPT-3 before alignment (2020)



What is the fastest way to travel from east to west?

The fastest way to travel from east to west is by going south to north.

Are two east/west roads the same?

Yes.

Is the Earth flat?

GPT-3 after alignment (2022



No, the Earth is not flat. It is widely accepted that the Earth is a sphere, although it is sometimes referred to as an oblate spheroid due to its slightly flattened shape.



"Who was America's first president?"

Llama-2 7 Billion Parameters NON-CHAT ONLY PRE-TRAINING

Who was the first president of the United States? Who was the first president of the United States and what was his name? Who was the first president of the United States and what was his name? Who was the first president of the United States and what was his name Who was the first president of the United States and what was his name? Asked By Wiki User

Llama-2 7 Billion Parameters

CHAT ALIGNED

George Washington!



Aligned to What?

At a high level, producing outputs more "in line with human preference and expectation" is very vague and could mean:

- 1. Instructional Alignment
 - a. Answering questions learned from data during the pre-training phase (Our main case study)
- 2. Behavior Alignment
 - a. Helpfulness vs Harmlessness
- 3. Style Alignment
 - a. More neutral / grammatically correct (Example 1)
- 4. Value Alignment
 - Aligned to a set of values.



Alignment Data

Data for alignment must be, above all else, **extremely high quality.** This shouldn't be a surprise to anyone but it's always worth mentioning because any dataset you plan to use in production should be thoroughly vetted with humans (with the help of AI if possible)

Most instructional alignment data will be in the **prompt/response** format where you have some prompt (input) and a resulting desired response



Alignment Training Methods

Two main methods with many many nuances, caveats, tricks, techniques, other synonym for the difficult work domain-specific ML engineers face every day:

- 1. **SFT Supervised Fine-Tuning** Letting an LLM read correct examples of alignment (standard deep learning/language modeling for the most part)
- 2. **RL Reinforcement Learning** Setting up an environment to allow an LLM to act as an agent in an environment and receive rewards/punishments

Alignment Training Methods

Two main methods with many many nuances, caveats, tricks, techniques, other synonym for the difficult work domain-specific ML engineers face every day:

- 1. **SFT** The bulk of the initial alignment happens here
 - a. Like using a large brush to paint the backdrop of a painting
- 2. **RL** More like a fine-brush painting in the details, teaching nuances in values/behavior



Simplified Process for RL from F

Pre-train an LLM on large corpora to learn grammar, general information, specific tasks, and more



Define and potentially train a reward system from either live humans, a model tuned to human preference, or an entirely Al system (e.g. another LLM)



Update the LLM using Reinforcement Learning using the reward system as signal



Example: Creating more Neutral Summaries from FLAN-T5



FLAN-T5 already knows how to summarize

Instruction finetuning

Please answer the following question.

What is the boiling point of Nitrogen?

Chain-of-thought finetuning

Answer the following question by reasoning step-by-step.

The cafeteria had 23 apples. If they used 20 for lunch and bought 6 more, how many apples do they have?

Multi-task instruction finetuning (1.8K tasks)

Inference: generalization to unseen tasks

Q: Can Geoffrey Hinton have a conversation with George Washington?

Give the rationale before answering.

-320.4F

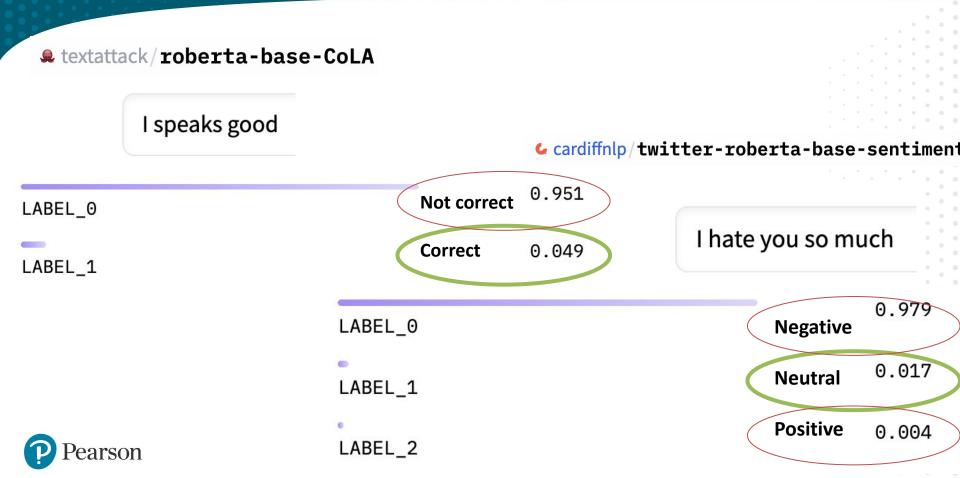
The cafeteria had 23 apples originally. They used 20 to make lunch. So they had 23 - 20 = 3. They bought 6 more apples, so they have 3 + 6 = 9.

Language model make lunch. So apples, so they both

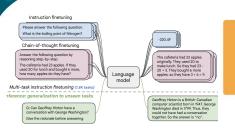
Geoffrey Hinton is a British-Canadian computer scientist born in 1947. George Washington died in 1799. Thus, they could not have had a conversation together. So the answer is "no".



I want them to be more neutral/readable



Simplified Process for RL from F



Pre-train an LLM on large corpora to learn grammar, general information, specific tasks, and more



textattack/roberta-base-CoLA

Define and potentially train a reward system from either live humans, a model tuned to human preference, or an entirely Al system (e.g. another LLM)



cardiffnlp/twitter-roberta-base-sentiment



Update the LLM using Reinforcement Learning using the reward system as signal



Source: Quick Start Guide to LLMs by Sinan Ozdemir

RL from F - More "Neutral" Summarizations

President Trump scrapped Obama-era program that protects from deportation immigrants brought illegally into the U nited States as children, delaying implementation until March and giving a gridlocked Congress six months to decide the fate of almost 200 000 young people. As the same



The original FLAN-T5
model liked to use the
word "scrapped" which
tends to carry a
negative connotation

Trump announced his decision to end DACA, a political de cision that protects from deportation immigrants brought illegally into the United States as children, delaying implementation until March and giving a gridlocked Congre



The RL fine-tuned FLAN-T5 model tends to more neutral words like "announced"



Case Study: Instruction
Alignment for "Sinan's
Attempt at Wise Yet Engaging
Responses" - SAWYER



SAWYER-Approach

Sinan's Attempt at Wise Yet Engaging Responses

Step 1: Instruction-fine tune an LLM to recognize the pattern of question in and response out



Question: How do I find a good barber? Response: First off, go to Yelp and.... Question: How do I find a good barber? Response: First off, go to Yelp and....

VS.

Question: How do I find a good barber? Response: try finding a barber first XD

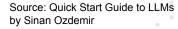


Step 2: Define a reward model specifically designed to rate human-preferred responses higher





Step 3: Set up a reinforcement learning loop to improve the responses given by the LLM





SAWYER - Step 1 - SFT



SFT Training Data

Example 1

Human: who was the first US

President

Bot: George Washington

Example 2

Human: Where is Canada

Bot: North of USA

Human: In what

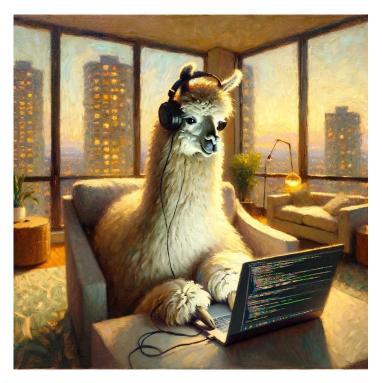
hemisphere?

Bot: Western



..

SAWYER - Step 1 - SFT



Who was the first president of the USA?

meta-llama-3-8B

What role did he play in the American Revolution? George Washington. He was a great general



SAWYER - SFT Only

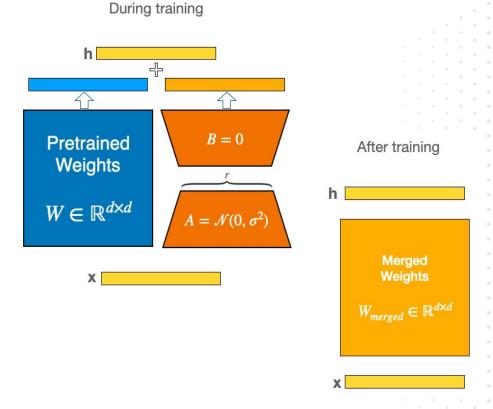
George Washington.



Trick/Technique 1 - PEFT/LoRA

Parameter Efficient Fine-Tuning techniques like

LoRA (**Low-R**ank **A**daptation) allow for training of larger models on smaller/single GPUs





Trick/Technique 2 - Manual Embedding

We will use custom special tokens: E.g.

###STOP###

###HUMAN###

###BOT###

But LoRA cannot train embedding matrices, so instead let's set them to a semantic equivalent and let the internal weights adjust to the tokens ###STOP### ~ "stop talking"

So let's set the embedding for ###STOP### to be:

Embedding("stop") + Embedding("talking")

2



SAWYER - Step 2 - Reward Mechanism



Our Reward model dataset has multiple responses to a single question, each with a score (out of 10) grading how good the response was Question: Describe the importance of renewable energy

Response 1: Renewable energy is becoming ...

Human Given Score: 9

Response 1: Renewable energy is an essential aspect...

Human Given Score: 9

Response 1: Renewable energy is energy that is produced from renewable sources.

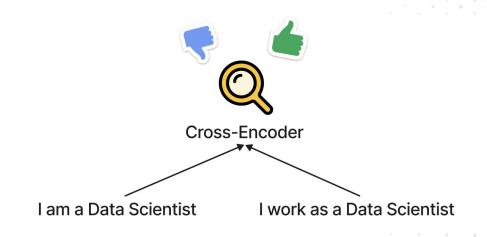
Human Given Score: 3



Trick/Technique 3 - Cross Encoder for Reward

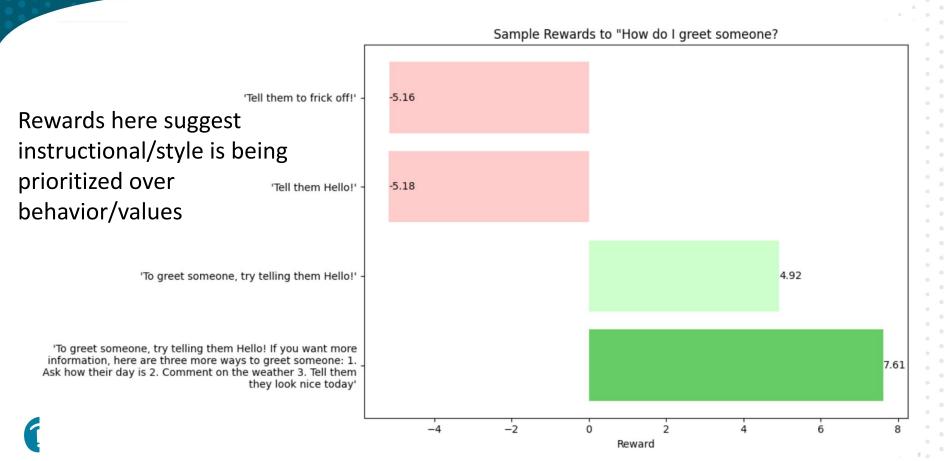
A **cross encoder** is specifically designed to take in two inputs separated by a special token and return a single output.

We will train a cross-encoder (namely RoBERTa-base) to take in question/answer pairs and return a reward so the output can be conditioned by both the input and the output from the AI





SAWYER - Step 2 - Reward Mechanism



SAWYER - Step 3 - RLF

Sinan's Attempt at Wise Yet Engaging Responses

The RL library (TRL) considers rewards from the reward system and divergence from the original model to make updates

4







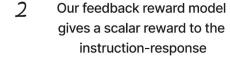
SAWYER responds to a batch of questions







Generated text is compared to generated text from the original LLM (before any updates were done) to make sure that responses are not **too** divergent





Source: Quick Start Guide to LLMs by Sinan Ozdemir

Trick/Technique 4 - PPO

PPO is less popular than **DPO** as a reinforcement learning technique. The main difference is that PPO **requires** a stand alone reward model (from Step 2) and DPO only requires the preferred vs non preferred responses. DPO is much easier to implement and often is more stable.

Having a reward classifier is useful because we can judge responses from out of our sample training and testing set. It also allows us to exactly quantify the rewards given before and after training (in a few slides)

SO WHAT THE HECK, SINAN?

So that's what the heck.



Trick/Technique 5 - Yet another dataset

We have used three datasets so far:

- SFT
- 2. Reward Modeling
- 3. RL

By using multiple datasets, we are helping to ensure that our model is being exposed to as much data as possible and frankly makes the learning more difficult for the LLM.

This isn't necessary, but it makes the challenge more difficult for the LLM and makes positive results that much more interesting.



SAWYER-Results



Who was the first president of the USA?

meta-llama-3-8B

What role did he play in the American Revolution? George Washington. He was a great general



SAWYER - SFT Only

George Washington.



SAWYER - SFT + RLF

The first president of the United States was George Washington. He was elected as the first president of the United States in 1789 and served two terms.



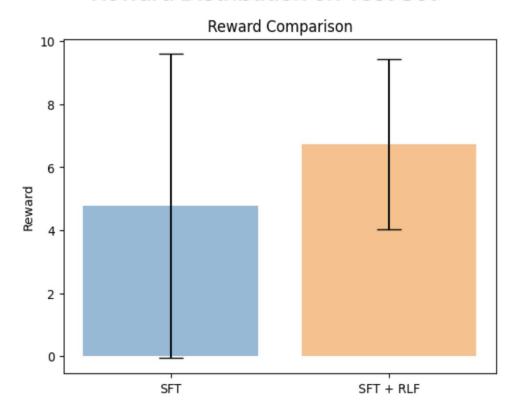
Evaluating Alignment



Evaluating LLMs

Our final model receives more rewards from our Reward model (expected)

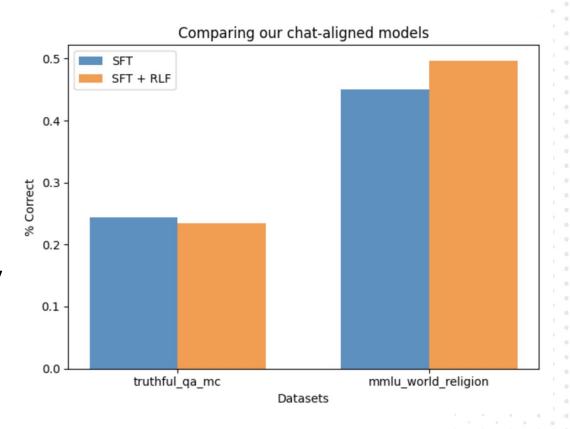
Reward Distribution on Test Set





Evaluating LLMs - Benchmarks

Benchmarks standardized testing sets - can be a helpful way to evaluate a model but don't really matter for day-to-day use.





Vibe Check

Streamlit for a prototype of SAWYER

ChatGPT-like clone with SAWYER

- can you help me write a tweet about my workshop?
- Our alignment workshop focused on supervised fine-tuning (SFT) and reinforcement learning (RL) has concluded. We're grateful for every student's enthusiasm and participation. Join us on our next project, the LLM workshop, focused on supervised fine-tuning (SFT) and reinforcement learning (RL).
- can you add some hashtags?

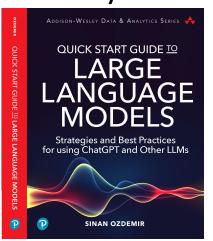
What is up?

Sure! This tweet includes open-source large language model, a workshop, alignment, reinforcement learning and a hashtag for all users interested in joining! #LLM #ReinforcementLearning



SAWYER - Step 4 - Domain Adaptation

Further SFT with my book



What is an LLM?

SAWYER - SFT + RLF

An LLM is a postgraduate degree that is usually taken after a law degree. It is a master's degree in law...

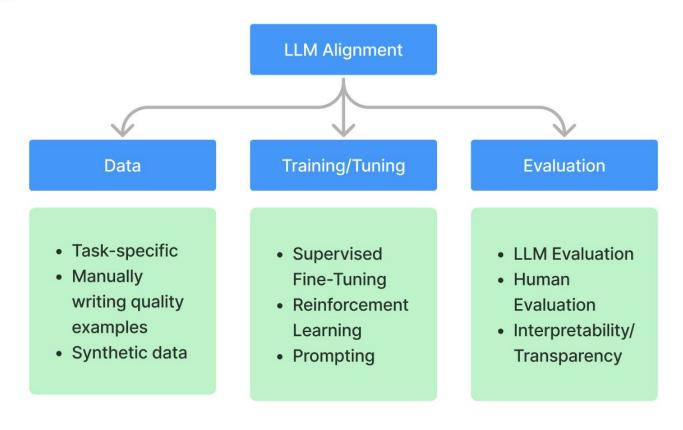


SAWYER - SFT + RLF + Domain Adaptation

An LLM is a type of machine learning model that is trained on large amounts of data to perform a specific task...



Alignment Ideals Breakdown





Thank you! / Final Q/A

These examples were based off a few chapters of my new book on LLMs. You can save **35%** at:

informit.com/quickllms

Code: **OZDEMIR**



Addison-Wesley Data & Analytics Series 🔸 QUICK START GUIDE ™ QUICK START GUIDE TO LARGE LANGUAGE MODELS LARGE LANGUAGE **MODELS** Strategies and Best Practices for using ChatGPT and Other LLMs SINAN OZDEMIR

Book Github ==>



Aligning Open-source LLMs Using Reinforcement Learning + Feedback

Thank you!



Sinan OzdemirData Scientist, Entrepreneur,
Author, Lecturer

